



The Sovereign Compute Access Act

Title I — Sovereign Compute Access. Reference standard, Sections 101–105 + proposed 106.

I	Sovereign Compute Access	101–106
II	Human Agency & Privacy	201–204
III	Federated Jurisdiction & Anti-Capture	301–305

*“Americans should not be required to rent
intelligence they can lawfully possess and operate
for themselves.”*



One tier is mandatory. The rest stay commercial.

Every Personal, Cooperative and Small Business account gets free, unmetered, private access to any open-weight model a provider publishes.

Tier I — Personal, Cooperative & Small Business

FREE ·
MANDATORY

Bank of Me and Bank of Us accounts, plus small business, get unmetered access to every open-weight model on the registry, run inside their own Sovereign Pod. No telemetry requirement. No API lock-in.

Tiers II & III — Pooled & Frontier

OPT-IN ·
PAID

Anything beyond local hardware — pooled GPUs, hosted inference, frontier-scale models — stays a commercial service, priced by each vendor. The Act does not force these open.

The Act asks nothing new to be opened. It asks that what is already published be guaranteed.



Six principles hold the registry together.

- 01 Capability, authority and compute footprint — never brand**
- 02 Total and activated parameters recorded separately**
- 03 Cross-ownership disclosed, not penalised** 103(a)
- 04 Etched weights committed before tape-out** Proposed 106(a)
- 05 Scheme, parameters and compiler published together** Proposed 106©
- 06 Inclusion is informative, not an endorsement**

A stated intention to release is not a release.



I Personal & Cooperative

BOM and BOU accounts, plus small business — unmetered, inside your own pod.

HARDWARE CLASS

Laptop, desktop, edge, phone

Runs inside a private M5POD, fully under the holder's control

PARAMETER BAND

0.3B – 32B

Dense or quantized; MoE evaluated by activated params

ALIBABA Qwen3-0.6B edge

ALIBABA Qwen3-4B / 8B personal

ALIBABA Qwen3-14B / 32B small business

OPENAI gpt-oss-20b 3.6B active, 131K ctx

GOOGLE Gemma 3 1B-12B

META Llama 4 Scout 17B active, 10M ctx

META Muse Glimmer single GPU / Mac, Aug 2026

MICROSOFT Phi-4 14B, MIT

APPLE OpenELM 270M-3B

DEEPSEEK R1-Distill-Qwen 1.5B-32B

Llama 4 Scout ships under the Llama Community License rather than Apache or MIT — flagged, but still open-weight for Tier I.



II Scaled & Pooled

Where access moves from owned hardware to paid pooling.

HARDWARE CLASS

Shared or rented GPU pools

Single or clustered datacenter GPUs, compute-as-a-service

PARAMETER BAND

30B – 500B

Total parameters; activation and context scale with pooled compute

ALIBABA **Qwen3-235B-A22B** 22B active

OPENAI **gpt-oss-120b** 5.1B active, single 80GB GPU

MISTRAL **Small 3.1** 24B, text+image, 128K ctx

GOOGLE **Gemma 3** 27B

META **Llama 4 Maverick** ~400B total, 1M ctx

DEEPSEEK **V4-Flash** 284B total / 13B active, MIT

NVIDIA **Nemotron 3 Super** 120B total / 12B active

MINIMAX **M3** 428B total / 23B active, 1M ctx

Pooling must stay interoperable and non-exclusive to any one chip vendor.

Nvidia is now a model publisher in its own right — a Section 103(a) disclosure relationship.





III Frontier & Spatial

Thinking labs, world models,
server-class research.

HARDWARE CLASS

Server-class, multi- node clusters

Beyond a single personal-device or single-GPU target

PARAMETER BAND

500B+ total

Spatial and world-model class
classified on five axes

DEEPSEEK **R1** 671B total / 37B active

DEEPSEEK **V4 Pro** 1.6T total / 49B active, MIT

Z.AI **GLM-5.2** 753B total / 40B active, MIT

MOONSHOT **Kimi K2.6** 1T total / 32B active

THINKING MACHINES **Inkling** 975B total / 41B active

NVIDIA **Nemotron 3 Ultra** 550B, Open MDW

PERPLEXITY **R1-1776** DeepSeek R1 fine-tune

META **V-JEPA 2** 1.2B, world model

XAI **Grok-1** 314B, unmaintained since Mar 2024

*Section 105(d) classifies spatial models on FLOPs,
coherence horizon, action space, modality fusion and
real-time constraint — not parameter count alone.*



Where the open weights are not.

Disclosed gaps, plus commitments not yet met.

Anthropic NO OPEN WEIGHTS

No open-weight model released to date. Claude remains closed-weight only.

World Labs - Google **DeepMind**

NO OR RESTRICTED
RELEASE

World Labs has released no open-weight spatial model; Genie 3 is closed and restricted to a research-preview group. Meta is the counterexample — V-JEPA 2 ships downloadable weights.

Meta Superintelligence Labs - AMI Labs

COMMITTED, NOT YET
SHIPPED

Muse Glimmer shipped open-weight Aug 10 2026, with Muse Spark 1.2 weights stated to follow. LeCun's AMI Labs (Paris, Mar 2026) has released nothing to date.

Inclusion in a tier is informative, not an endorsement.



Cross-ownership is disclosed, not penalised.

**Who holds equity in whom,
stated plainly on the record.**

Nvidia CROSS-EQUITY DISCLOSED

H100, H200, Blackwell B100/B200. Disclosed equity or committed investment in OpenAI, xAI, Anthropic and Intel, plus the Groq asset purchase — the dominant vendor across nearly every major lab.

AMD MODEL-SPECIFIC SILICON

Instinct MI300X and MI325X, independent on model-publisher equity. Signed to acquire Taalas Aug 6 2026, closing Q4 — bringing etched-weight parts into scope for Proposed Section 106(a).

Intel CROSS-EQUITY DISCLOSED

Gaudi 3, Falcon Shores and the Heracles FHE accelerator. Has accepted an Nvidia investment commitment — itself disclosable as a chip-to-chip relationship.



When the weights are the chip.

An etched part cannot be asked what it runs. The record is made before the mask.

Taalas HC1 - AMD WEIGHTS ETCHED IN MASK ROM

815mm² die on TSMC 6nm, 53B transistors, Llama 3.1 8B at 16,960 tokens/sec with no DRAM read per forward pass. One model per part.

Etched Sohu - OpenAI Jalapeño FIXED-FUNCTION INFERENCE

Hard-code a narrower slice of the workload into silicon for throughput a general-purpose part cannot reach.

Neuromorphic and analog in-memory COMPUTE-NATIVE WEIGHTS

Spiking and compute-in-memory parts hold weights in the device itself. The obligation is identical: the weights are physical.

Publish the weight digest, model version and mask-set identifier before the mask is cut, or the part is unverifiable for life.



The data is never decrypted.

Purpose-built parts took FHE from research to service.

Intel Heracles FHE ACCELERATOR · 3NM

ISSCC 2026, from a five-year DARPA programme. 8192-way SIMD, 64 tile-pairs, 48GB on-package memory. 1,074×–5,547× a 24-core Xeon on FHE math; a voter-privacy check in 14µs against 15ms.

Ant Misuan MPU H1 VENDOR CLAIM

BFV, BGV, CKKS and TFHE, including encrypted large-model inference. Claims over 10,000× CPU on core FHE operations.

FPGA, PIM and research ASICs RECONFIGURABLE AND IN-MEMORY

TFHE processors on AMD Virtex UltraScale+ HBM parts; CraterLake, SHARP and FHEmem set the published baselines.

*S4 and S5 stop being promises about handling and
become properties of the part.*



If the compiler is closed, the privacy claim is unverifiable.

Scheme, parameter set, compiler and version published together, or not claimed.

No proprietary compiler, runtime or attestation service may be the only path to a compliant deployment.

GOOGLE **HEIR** MLIR compiler for FHE

ZAMA **Concrete** TFHE compiler

ZAMA **TFHE-rs** Rust TFHE library

OPENFHE **OpenFHE** BGV / BFV / CKKS / TFHE

MICROSOFT **SEAL** BFV / CKKS

EPFL **Lattigo** Go lattice library

HE.ORG **Security standard** parameter tables

MLIR **Portability** ASIC · FPGA · GPU · CPU

Hardware neutrality and toolchain neutrality are the same principle one layer apart.



Closed silicon, reopened by an open OS.

**Where the vendor documents nothing,
reverse engineering is the only route left.**

Asahi Linux · M1 and M2 DAILY DRIVER

Conformant OpenGL 4.6 and Vulkan 1.3, plus Wi-Fi, Bluetooth, audio, webcam and display. Fedora Asahi Remix is the reference distro.

M3, M4, M5 bring-up WORK IN PROGRESS

August 2026: M3 put weeks away after the 7.2 cycle fixed webcam, microphone, USB 3.0 and Thunderbolt — installer, GPU and display still WIP. M4 and M5 are early bring-up.

Omarchy · omarchy-mac ARCH + HYPRLAND

Omarchy publishes x86_64 only; aarch64 returns 404. Apple Silicon runs the omarchy-mac fork on Asahi Alarm with community aarch64 builds.

Every driver here exists because someone reverse-engineered an undocumented block.



Five commitments, one public record.

Each layer publishes a timestamped hash. A claim made later than its commitment is not a disclosure.

- | | | |
|-----------|---|----------------|
| L1 | Silicon — weight digest, model version, mask-set ID | at tape-out |
| L2 | Runtime — FHE scheme, parameters, compiler version | per deployment |
| L3 | Ownership — cross-equity and vendor relationships | 103(a) |
| L4 | Ledger — timestamped hash commitments, public | ICSN |
| L5 | Classification — account, asset and protection level | M5CANON |

M5 economics can be built at the chip layer only if the chip layer is on the record.



No new purchase required.

**Register the computer you already
own as a Registered Sovereign Node.**

Turn existing hardware into a node

ANY CHIPSET, ANY
MAKER

Tens of millions of functional computers already sit idle or under-used in American homes. Any Covered Tier I holder can register their own PC or Mac as part of a decentralized Internet of Sovereign Things.

Registration verifies capability, not content

OWNER
CONTROL

Validation confirms hardware capability and a tamper-evident baseline. It grants no party — not TitleChain Foundation, not ICSN, not any provider — access to prompts, models, files or usage.

M5-H0 → H5 classification

REGISTRATION
PATH

Produces a device-identity and capability record consistent with ICSN Hardware Onboarding Standard RFC 0001, with the node commitment written to L4.

Eligibility never depends on chipset or brand.



Public compute earns a public return.

**The incentive to open more models,
without compelling private work.**

If public compute was used

REQUIRED, NOT
OPTIONAL

A model trained with a government grant, subsidized national compute, preferential export treatment or a direct contract must have an open-weight variant released within 18 months. Modeled on Bayh-Dole.

If a company opens up voluntarily

COMPLIANCE GETS
CHEAPER

Registry recognized status: procurement consideration under 102(b)(4) and self-attestation in place of third-party evaluation under 102(b)(5) — a direct cost saving, not a badge.

*No provider is compelled to open a model built
entirely on private compute.*



Nobody should pay for a baseline that already exists in the open.

Every major publisher already ships something open. The Act guarantees it rather than leaving it to discretion.

M1	Bank of Me — free Tier I access	Dimension 1
M2	Bank of Us — free Tier I access	Dimension 1
S5	Sovereign protection, attested in silicon	Dimension 3

Every tier runs inside a private M5POD under the holder's own authority.

Read with Appendix B — one standard from two directions.

